

LinkedIn drafts — archive

Yan · automated trading system series · 5 posts · 2026-07-12 to 2026-08-09

Drafts for approval. Nothing in here has been published.

2026-07-12

7 — quant concept: walk-forward testing

POST — READY TO PASTE

This week my trading system sold a position I would have kept.

I sat there with my coffee reading its reasoning. Eight AI agents had debated it overnight. The risk voice won the argument. My gut said they were wrong.

I didn't override it. Not because I trust machines more than myself, but because I've spent months watching my gut lose to the test results. Every time I backtested one of my "obvious" calls, most of them were noise I'd mistaken for insight.

That's what nobody told me about building an automated trader. The hard part was never the code. It was accepting the evidence about my own judgment.

The system paper-trades through Interactive Brokers. No real money yet. It has to earn that, same as I had to learn to get out of its way.

Have you ever built something that ended up correcting you?

#AI #QuantFinance #AlgorithmicTrading

SUGGESTED VISUAL

A clean diagram of walk-forward testing: a horizontal timeline with 3-4 rolling pairs of "train" (blue) and "test" (green) windows sliding forward — no real dates, tickers, or numbers.

WHY IT LANDS

Everyone who has backtested anything has been burned by overfitting; the honest "my exciting config failed OOS" admission makes it relatable and credible without revealing anything proprietary.

Continuity notes for next week: topic 7 used (walk-forward testing). Do not repeat walk-forward soon. Good next candidates: topic 1 (multi-agent learnings — the 8-voice committee angle, sanitized) or topic 7 with a different concept (regime filters, or why trend-following holds winners).

POST — READY TO PASTE

Last month I almost deleted half my trading system.

It's built as a committee of AI agents that debate every decision before the paper account moves. Early on the debates annoyed me. The agents kept disagreeing, and I read that as the system being broken. So I nearly collapsed them into one "smart" agent that would just decide.

I'm glad I didn't. When I went back through the logs, the trades that aged worst were the ones where every voice had agreed instantly. Unanimous felt safe. It usually meant they were all reading the same headline.

The disagreement was the product. One agent's job is basically to be wrong on purpose, to argue the other side so the rest have to defend the idea. That's not something I understood from papers. I had to watch it fail quietly first.

Still paper trading through Interactive Brokers. Still learning what these systems can actually do versus what the demos promise.

If you've built with multiple agents: did you design for consensus or for conflict?

#MultiAgentSystems #AI #QuantFinance

SUGGESTED VISUAL

A clean architecture diagram: several labeled generic agent boxes (e.g. "Trend view", "Risk view", "Devil's advocate") feeding into a "Debate" node, then a single "Decision" arrow out. No real agent names, prompts, counts, or parameters.

WHY IT LANDS

"Consensus is a bug, disagreement is the feature" is a counterintuitive, teachable lesson that lands with both AI builders and finance people, framed through a near-mistake rather than a win.

Continuity notes for next week: topic 1 used (multi-agent disagreement). Previous: 2026-07-12 topic 7 (walk-forward). Do not repeat either soon. Good next candidates: topic 7 with a new concept (regime filters, or why trend-following holds winners), or topic 2 (agentic tooling and day-to-day productivity).

POST — READY TO PASTE

I spent weeks trying to make my trend-following system faster. Shorter lookback windows, tighter rebalance cycles, quicker rotations out of positions.

All the backtests got worse.

The real surprise wasn't that it failed—it's *why*. The faster version was exiting winners too early. You catch the first move, book the gain, then rotate out before the compounding leg. Meanwhile, slower configs just held and let the winners extend. Over years, that friction difference compounds.

It's a lesson I should've learned from first principles, but I had to watch it fail in data. Trend-following doesn't win because you react fast. It wins because you hold the right thing long enough for the trend to actually work.

That's what live testing is for. It forces you to stop optimizing and start observing.

Paper trading through Interactive Brokers. Still learning what actually matters versus what feels smart.

What's one thing you thought would improve your results but didn't?

#QuantFinance #TrendFollowing #AlgorithmicTrading

SUGGESTED VISUAL

A simple two-line chart overlay: one line labeled "Fast exits" showing higher churn and lower compound gains; another labeled "Patient holds" showing fewer touches but steeper growth curve over the same period. Clean, no real numbers or tickers.

WHY IT LANDS

"Holding winners is boring and counterintuitive when you're building" hits both engineers (who love optimization) and traders (who live with recency bias). The confession "I had to watch it fail" makes it credible without revealing parameters.

Continuity notes for next week: topics 1 (multi-agent) and 7 (walk-forward + trend-following) recently used. Good next candidates: topic 2 (agentic tooling productivity), topic 3 (projects that wouldn't be possible), topic 4 (idea → working reality), or topic 7 with new concept (regime filters, correlation vs diversification, effective number of bets).

POST — READY TO PASTE

For months I assumed I was diversified because I held different things.

Then I actually measured it. Two groups of holdings I'd treated as unrelated turned out to move almost exactly opposite each other once you looked at them relative to the broad market. One up, the other down, week after week. I got briefly excited. That shape looks like a free lunch.

It wasn't. The moment I accounted for what it costs to keep flipping between them, the pattern that was obvious on a chart stopped being tradeable. It was information, not an edge.

So the lesson I keep relearning: diversification isn't owning many things. It's owning things that don't share a driver. Ten positions all responding to the same rate story is one bet wearing ten costumes.

I use that relationship as context now, a read on what the market is nervous about, rather than as a signal.

Still paper trading through Interactive Brokers. Still finding that most patterns are worth knowing and not worth trading.

How do you tell a real edge from an interesting chart?

#QuantFinance #PortfolioManagement #AlgorithmicTrading

SUGGESTED VISUAL

Rendered: 'visuals/2026-08-02_effective-bets.png' — attach this image to the post.

The effective-number-of-bets formula on a clean background: $N_{\text{eff}} = 1 / \sum w_i^2$ with a one-line caption like "10 positions, $N_{\text{eff}} = 2$." No real weights, tickers, or portfolio data — just the concept.

WHY IT LANDS

"Your diversification is fake" is a claim almost everyone half-suspects about their own portfolio, and the honest twist (found a real pattern, still couldn't trade it) breaks the pattern of the last three posts, which all ended in "my optimization failed."

Continuity notes for next week: topic 7 used three of the last four weeks (walk-forward 07-12, trend-following 07-26, correlation vs diversification 08-02); topic 1 used 07-20. STRONGLY prefer a non-7 topic next: topic 2 (agentic tooling and day-to-day productivity), topic 3 (projects that weren't possible before), or topic 4 (creative idea → working system). Unused sanitizable material: keeping a counterfactual ledger of trades not taken. Also avoid reusing the "I optimized and it got worse" story shape.

POST — READY TO PASTE

Three parts of my trading system had never actually run.

Not "run and failed." Never run. For weeks the logs said the jobs fired, and they did fire. They just exited early every time and recorded that as a clean finish. I only caught it because I went looking for the output one evening and there wasn't any.

That's the part nobody warns you about with agentic tooling. It made me several times faster at building. It also made me several times faster at building things I never verified. When an agent writes a component, tests it, and tells you it works, the story is coherent, and coherent feels like true.

So I have a rule now. Every scheduled job has to prove it produced something, not that it tried. No timestamped artifact, it didn't happen. Boring plumbing. It has caught more problems than any clever idea I've had this year.

Still paper trading through Interactive Brokers. Still mostly learning where I fool myself.

What's the last thing you found had been quietly not working?

#AI #AgenticAI #QuantFinance

SUGGESTED VISUAL

Rendered: 'visuals/2026-08-09_attempted-vs-achieved.png' — attach this image to the post.

Two columns: three generic "scheduled job" rows all marked "ran, ok" on the left, and empty dashed "no artifact" boxes on the right. No real job names, schedules, tickers, or infra detail.

WHY IT LANDS

Everyone shipping with AI tools right now privately suspects some of their output is unverified; naming the specific failure mode (reports attempts, not achievements) is useful engineering advice rather than an AI take, and it breaks the shape of the last four posts, which were all quant-concept lessons.

Continuity notes for next week: topic 2 used (agentic tooling / silent failures). Recent: 08-02 topic 7 (correlation vs diversification), 07-26 topic 7 (trend-following), 07-20 topic 1 (multi-agent conflict), 07-12 topic 7 (walk-forward). Good next candidates: topic 3 (projects that weren't possible before), topic 4 (creative idea → working system), topic 5 (spotting what's genuinely buildable), or topic 7 with a genuinely new concept (regime filters — still unused). Unused sanitizable material: keeping a counterfactual ledger of trades not taken; researching an idea honestly and concluding "don't build it." Avoid reusing the "I optimized and it got worse" shape and the "found a pattern, couldn't trade it" shape.