

What we feed the agents, and what they do with it

One context blob, sliced per agent

agents/context.py:build_context(state_dir, data_dir) runs once per cycle and does no network I/O — it reads local state files and the parquet cache only. Everything every agent sees comes from here.

Key	Source	Meaning
account, positions	runner.db snapshot	the traded book — equity, per-name qty, 52w entry/now percentile, unrealized, sector
book_concentration	computed	effective number of bets (ENB) from correlation eigenvalues
candidate_ladder	agents/candidates.py	the ranked scoreboard — the only source of names not already held
holds, k_override	state/holds.json	operator pins
operator_objections	copilot thread	what Yan argued recently, in his words
operator_mandates	copilot/mandates.py	standing instructions, graded strong/medium/soft
macro_dial	macro_monitor.json	macro readings
vol_surface	options_monitor.json	implied/realized vol
spy_vix_triggers, style_leader	advisor.json, style_advisor.json	
established_lessons	memory	up to 6 — what the desk has learned
operator_lessons_under_consideration	memory	Yan's un-hardened lessons
dossiers, source_trust	memory	world-state notes, per-source Beta trust
economy	FRED via econ_watch	CPI, claims, HY spreads
sector_momentum_vs_spy_pct, headlines	news_watch	last, so it is first to be trimmed

_VIEW_KEYS in committee.py then hands each persona only its slice. The creative agent is deliberately denied positions — position blindness is the whole point of that role.

The eight personas

Agent	Sees	Job
quant	ladder, positions, macro, vol, economy	Reason from numbers. Two hard rules: a high 52w percentile says where price is, not that reward-to-risk is good; correlated holdings are ONE bet.
narrator	macro, dossiers, trust, lessons, headlines	Geopolitics, CB psychology, what is priced in
street	ladder, positions, style, sector momentum, headlines	Sell-side ratings, revisions, crowded love/hate
position_coach	positions, holds, lessons, objections, mandates	Only our book — does each thesis still hold
risk_officer	positions, vol, macro, triggers, economy	What can hurt us this week. May say "nothing actionable".
trader	ladder, positions, vol, triggers, headlines	Tactical, this-week horizon
scout	ladder, sector momentum, headlines, trust, lessons	The NEXT theme, weeks to months. Carries a standing operator directive on quantum computing.
creative	ladder, sector momentum, headlines, economy, macro, trust — not the book	The idea the committee is missing because they are staring at their holdings. Weighs named social sources by trust score.

Each returns the same schema:

```
{
  "stance": "bullish|neutral|bearish",
  "take": "<4-6 sentences, specific>",
  "prediction": {
    "subject": "NVDA",
    "direction": "up",
    "horizon_days": 14,
    "confidence": 0.72
  },
  "sources": ["..."],
  "cited_lessons": ["ls-..."]}

```

The prediction is the point. It is stored, and graded weeks later against real closes, producing the per-agent hit rate and Brier score that everything downstream uses to weight the voice. Track record, not eloquence.

The debate

1. Specialists (mid-tier model) → 8 takes, each journaled as kind="take", each prediction stored.
2. Guards — deterministic, non-LLM checks the personas miss. Advisory; never an order gate.
3. Challenger (frontier) — reads all takes plus market context, attacks the two highest-conviction ones.
4. Manager (frontier) — synthesises posture from takes, objections, calibration, guard flags, and the lesson book. Journaled as kind="committee" — note this row stores only the ruling, not the takes.

The PM (weekly, frontier model)

Reads, in prompt order:

```
operator_mandates, operator_held_do_not_trade,
sim_portfolio { equity, cash, holdings, performance, cycles_since_name_change },
allowed_universe_etfs (21 ETFs), allowed_universe_stocks (~1600, as a note),
agent_takes (16 most recent),
agent_calibration,
sentinel_alert,
today_context ← with live-account keys renamed *_NOT_YOUR_BOOK
recent_committee_rulings (6)
```

Assembled by `_budgeted_prompt`, which fits 24k characters by dropping whole items — headlines first, then oldest history — never by slicing the string. The book, the caps, the mandates and today's context survive every reduction, so the result is always valid JSON.

Returns target weights, a rationale, and a watch item. Weights then pass through `_clamp_weights`: whitelist → per-name cap → ETF-count cap → cluster cap → gross cap. Anything off-whitelist or operator-pinned goes to cash, never to a creative substitute.

The PM cannot construct an **Order**. It writes `state/agent_pm/` and sends Telegram. That isolation is structural, not conventional.

The historian (Fridays)

Reads a week of journal plus the current lesson book, proposes at most two new candidate lessons, and votes on existing ones. Promotion to established needs +3 net supporting evidence.

The copilot (on demand, Telegram)

Answers questions over `NOW_trading_account_paper`, `NOW_trading_account_orders`, `NOW_agent_pm_simulated_book`, `NOW_risk_state`, `NOW_lesson_book`, plus recent chat turns. Read-only by construction — `tests/copilot/` asserts the package never imports `execution/`. It can capture mandates and, since 2026-08-06, discuss the lesson book.

Feedback loops — which are closed

Loop	Closed?
prediction → graded vs price → agent calibration → weighting	yes, nightly
source cited → prediction graded → source trust	yes
PM decision → daily mark → equity curve vs SPY	yes
candidate considered but passed → shadow book → forward return	yes
lesson → cited by an agent → credited	no — cited_lessons is collected and never read
closed trade → episode → lesson evidence	no — add_episode is never called

The last two are the subject of finding #1 in [04_PRODUCTION_READINESS.md](#).