First-month live checklist

Print this. Tick it. The point of the first month is not profit — it is finding out what is broken while the amount at risk is small enough not to matter.

Everything on this list exists because something in this system once ran without erroring, without alerting, and without working. On 2026-08-06 a single sweep found the agent scorecard had never graded a prediction, the counterfactual ledger had never graded a leg, and the historian had never been scheduled on a box without an unrelated feature flag set. All three were green on every check that existed at the time.

So the governing rule of this checklist: never accept ”it ran” as evidence. Ask what it produced.

Before you switch it on — day 0 gate

Do not proceed past a single unticked box.

☐	Check	How	Pass
☐	Account holds only the money you can lose	IBKR portal	8–10k, everything else moved out
☐	Both live gates deliberately set	`grep -E "TRADING_ENV|ALLOW_LIVE" .env`	`live` and `true`, typed by you today
☐	Risk limits reviewed line by line	`uv run trading status`	position/gross/daily-loss/drawdown are what you intend for a 10k account
☐	Margin borrowing off	`.env`	`MAX_MARGIN_BORROWING_PCT=0.0`
☐	Cycle approval on for month one	`.env`	`REQUIRE_CYCLE_APPROVAL=true` — you approve every rebalance by hand
☐	IBKR 2FA resolved for unattended login	test a gateway restart	comes back without you tapping a phone
☐	Full suite green on the Mac	`uv run pytest -q -m "not slow and not live"`	0 failures
☐	VPS has headroom	`free -m`	≥800 MB available, swap < 20%
☐	Deployed image is the commit you think	`git log --oneline -1` on VPS + `docker compose ps`	container ages in minutes, not hours
☐	Kill switch rehearsed	`/halt` then `/resume` in Telegram	halt takes effect and is visible in `/status`
☐	You know how to flatten manually	IBKR portal, not the bot	you have done it once
☐	Baseline snapshot recorded	see ”Week 0 baseline” below	written down on paper

Open items already logged in **`TODO.md`** / **`GO_LIVE.md`**

Found 2026-08-06 by re-reading the backlog. These predate this document and several are live-blocking. Do not treat the day-0 gate as complete until each has been resolved or consciously accepted.

☐	Item	Why it matters live	Source
☐	Order status never promoted for overnight fills	An after-hours order that fills at the next open stays `submitted` in `orders.db` forever. The state machine loses track of a real position. On live that risks re-submitting or mis-sizing.	TODO Phase 11

☐	Item	Why it matters live	Source
☐	No single execution lock	Cron cycle, trigger cycle, approval flow and manual commands have no shared mutual-exclusion. Per-job locks + a 10s cooldown mitigate; they do not prevent. Order stacking already shorted two names on paper once.	TODO Phase 11, external review 2026-07-15
☐	**`trader-live`** compose service is stale	Shares the paper state dir and uses bridged networking. Live must have a fresh state dir — sharing halt state, order ledger and snapshots between paper and live is a data-corruption path.	TODO Phase 10
☐	Sized-down live limits not yet set	Plan on record: `MAX_POSITION_PCT=0.05`, `MAX_GROSS_EXPOSURE=0.50` for month one, kills unchanged.	GO_LIVE §3
☐	**`live-candidate-1`** tag after one clean cycle on the deployed commit	You cannot roll back to "the version that worked" without one.	GO_LIVE §2
☐	IB Gateway live login — live username is the paper username without the trailing "2"; port 4001 not 4002; needs `up -d --force-recreate`, a plain restart does not re-read `.env`	Getting this wrong on live day means no connection at the open.	GO_LIVE §3
☐	Abort procedure written on paper	How to halt, how to flatten manually in TWS, what happens if the VPS dies mid-session. Literally on paper.	GO_LIVE §3
☐	Telegram alerts on a channel you watch with sound on	First live week.	GO_LIVE §3
☐	"No trades in N cycles" watchdog — still open	The backlog says it "would have caught the June dead month in week one. Small; do first." The new ops liveness checks cover the committee/PM/historian, not "the cycle produced no orders".	TODO Phase 11
☐	Winter DST — before November set `CRON=5 22 * * FRI`	21:05 UTC is only 5 minutes after the close in winter.	TODO Phase 10
☐	Execution quality — raw market orders today	Adaptive algo or marketable-limit protects against bad opens. Not blocking at this size, but it is real money now.	TODO Phase 11

Week 0 baseline — write these numbers down

You cannot detect drift without a starting point.

```
cd /opt/trading-agent && docker compose exec -T trader python - <<'EOF'
from trading.memory.store import default_store
m = default_store(); s = m.stats()
print("journal", s["journal"], "| episodes", s["episodes"],
      "| lessons", s["lessons"], "| predictions", s["predictions"])
print("calibration:", m.calibration())
EOF
```

Metric	Day 0 value
account equity	
open positions	
predictions recorded / graded	
episodes	
lessons established / candidate	

Metric	Day 0 value
agents with a calibration row	

Every trading day — 5 minutes

1. Features working

☐	Check	Pass
☐	Ops channel silent, or every alert understood	silence = healthy; an alert you cannot explain is an incident
☐	Daily summary arrived (20:10 UTC)	present, and the P&L matches the dashboard
☐	Cycle ran and you approved or rejected it deliberately	no cycle auto-expired unattended
☐	`/health` and `/positions` respond	bot alive

2. Nothing stale

☐	Check	Pass
☐	Dashboard "SYSTEM" panel ages	broker snapshot < 1h, news < 24h, macro < 24h
☐	Price cache reached yesterday's close	ladder `as_of` is the last trading day
☐	No "[WARNING] price cache still stale after refresh" alert	absent

3. No abnormal behaviour — read every fill

This is the one that catches the thing nobody predicted. Look at each trade and say out loud why it happened.

☐	Check	Red flag
☐	Every fill maps to a strategy signal you can name	a name you do not recognise
☐	No position exceeds `MAX_POSITION_PCT` of equity	any breach = stop and investigate
☐	No short positions	system is long-only; a negative quantity is a bug
☐	No duplicate/stacked orders in the same name	order stacking shorted two names on paper once before
☐	Fill prices near the close you expected	> 1% slippage on a liquid name needs an explanation
☐	Currency: no unexpected USD cash negative	margin auto-loan starting

---

Every week — 30 minutes, Saturday

4. Is it learning?

The whole reason for the first month. Run this and compare to last week:

```
cd /opt/trading-agent && docker compose exec -T trader python - <<'EOF'
from trading.memory.store import default_store
m = default_store(); s = m.stats()
print("predictions:", s["predictions"], "| episodes:", s["episodes"], "| lessons:", s["lessons"])
print("\ncalibration (agent, n, hit rate, brier):")
for c in m.calibration(): print(" ", c)
print("\nestablished lessons:")
for r in m.lessons(status="established")[:10]: print(" -", r["statement"][:160].replace("\n"," | "))
EOF
```

☐	Check	Pass	Fail means
☐	`predictions graded` increased	strictly higher than last week	the grader is dead again
☐	`calibration` has a row per active agent	8 rows eventually	predictions recorded but not scored
☐	Brier scores are moving	not frozen	grading stuck
☐	`episodes` increased if anything closed	matches your round-trips	`add_episode` not firing
☐	Historian ran Friday	a `historian` journal row dated this week	never-scheduled bug returned
☐	New lessons carry applies_when / fails_when / retire if / sample	all four present	historian is proposing superstitions
☐	`/edge` returns numbers	picks vs passes populated	shadow ledger not grading
☐	No lesson established on < 20 distinct names	check the `sample:` line	over-fitting

The trap: a lesson that reads as obviously true. Your own buy-quality-after-drawdowns idea tested at +5.8% pooled over eight years and below 50% hit rate every year since 2023. If a new lesson has no stated failure mode, it has not been thought through.

5. Configs respected

☐	Check	How
☐	Effective limits match `.env`	`uv run trading status` — remember **`risk.yaml`** portfolio numbers are NOT read by the risk manager
☐	No `/hold` pin silently dropped	`/holds`
☐	Mandates expired as expected	`/mandates`
☐	Universe refreshed Sunday	membership-delta alert arrived, or ”no changes” in logs
☐	Model routing unchanged	`AGENTS_MODEL*` unchanged unless you changed it
☐	API credit remaining	a zero balance returns HTTP 400 and silently kills the committee

6. Pipelines end to end

Trace one real decision all the way through. Pick a fill from this week:

☐	Stage	Where to look
☐	signal	`/signal` — was the name on the ladder?
☐	selection	cycle report — did it make the top-K?

☐	Stage	Where to look
☐	risk decision	order store — accepted, scaled, or blocked, and why
☐	order	`/orders` — one order, correct side and size
☐	fill	`/positions` — quantity matches
☐	snapshot	dashboard equity moved by roughly the right amount
☐	memory	an `episode` when it closes

If any stage cannot be traced, that is the finding. Write it down.

Every month

☐	Check	Pass
☐	Realised vs expected slippage	within the 5bps model, or update the model
☐	Live vs paper divergence	same strategy, same universe — large gaps mean an execution problem, not an alpha problem
☐	Drawdown vs `MAX_DRAWDOWN_PCT`	never approached without the halt firing
☐	Restore drill	`docs/restore.md` — restore state to a scratch dir and confirm it loads
☐	Disk trend	`df -h` — Parquet and logs grow forever
☐	Read `docs/incidents.md`	every incident this month written down

Incident rules — agree these with yourself now, not at the time

Trigger	Action
A trade you cannot explain	`/halt` immediately, investigate before resuming
A position over its cap	`/halt`, flatten manually via IBKR
Any short position	`/halt` — long-only invariant broken
Gateway down > 1 hour in market hours	flatten manually if holding, do not wait
Two consecutive days with `graded_today: 0`	learning loop dead — not urgent for money, urgent for the experiment
Ops alert you do not understand	treat as an incident until proven otherwise

Halting is free. Being wrong about why a trade happened is not.

What ”success” looks like after one month

Not profit. Profit over one month on 10k is noise, and treating it as a verdict is the single most likely way to fool yourself.

Success is:

- every trade explainable,
- no cap breached, no invariant violated,
- prediction count graded rising every week,
- at least one lesson established with conditions, or an honest none,
- at least one bug found — if you found nothing, you were not looking,
- and the incidents log has entries.

A clean month with an empty incidents log usually means the monitoring is broken, not that the system is perfect. That is the lesson of 2026-08-06.